Investment Memo: AgentVC Date: August 7, 2025 **Stage:** Pre-seed **Recommendation: WATCHLIST** - The company is tackling a high-value problem with a clear and compelling vision, but it is too early to invest given the lack of founder background information and nascent product traction.

---

1. Executive Summary (The "TL;DR")

- **The Core Thesis:** AgentVC is building a B2B SaaS platform that uses agentic AI to automate the entire top-of-funnel workflow for Venture Capitalists, from deal sourcing and screening to analysis. The thesis is that by creating a unified, intelligent, end-to-end system, they can solve the immense inefficiency VCs face when stitching together manual processes and fragmented point solutions, allowing them to focus only on the most promising deals.
- **Key Strengths:**
    1. **High-Value Problem:** The team correctly identifies a significant "painkiller" problem for VCs: the overwhelming and costly manual effort required for deal sourcing and screening.
    2. **Strong Strategic Positioning:** AgentVC clearly differentiates itself from legacy data providers (e.g., PitchBook), integrated CRMs (e.g., Affinity), and single-point AI tools by offering a comprehensive, workflow-native platform at an accessible price point.
    3. **Clear & Logical GTM:** Their go-to-market strategy of leveraging a Design Partner Program to build case studies before converting pilot users to paying customers is a proven, capital-efficient approach for enterprise SaaS.
- **Critical Weaknesses & Risks:**
    1. **Unknown Team:** The deck provides no founder bios or background information, making it impossible to assess founder-market fit, technical capability, or past execution history. This is a critical information gap.
    2. **Nascent Traction:** Progress is limited to 10+ VC interviews and one firm committed to a design partnership. There is no MVP, no usage data, and no revenue, making this a highly speculative, pre-traction investment.
    3. **Weak Initial Moat:** The primary defensibility mentioned is a "Proprietary Data Flywheel," which is a future goal, not a current asset. The company currently lacks a strong, defensible moat against fast-following competitors or incumbents.
- **The Ask:** The deck indicates they are preparing for a Pre-Seed fundraising round in Q4 2025. The specific amount and valuation are not provided.

2. The Founding Team (Score: 2/10)

- **Founder-Market Fit: (Critical Information Gap)** The pitch deck lists Rohan Vyas as CEO and Dhruv Puri as CTO but provides no biographical information, work history, or educational background for either founder. It is impossible to analyze their domain experience, unique insights, or any "unfair advantage" they may possess. This is the single biggest red flag in the deck.

- **Team Completeness:** The team consists of a CEO and a CTO, a standard co-founding duo. However, for their next stage of growth post-launch, they will need to hire key Go-to-Market personnel, as outlined in their roadmap for Q2 2026.
- **Execution Capability: (Critical Information Gap)** Without founder backgrounds, there is no evidence of past execution, grit, or ability to ship a product. We cannot analyze past projects, career progression, or previous entrepreneurial ventures.
- **Red Flags:** The complete omission of founder backgrounds in a pitch deck is a major red flag. It prevents due diligence on the most critical factor of a pre-seed investment and may suggest a lack of relevant experience.

3. Problem & Opportunity (Score: 9/10)

- **Problem Significance:** The deck excellently articulates a "painkiller" problem. VCs and accelerators are inundated with thousands of applications and spend excessive manual effort on sourcing and screening. The current solutions—either laborious DIY sourcing or expensive, often incomplete data subscriptions—are highly inefficient. The quotes from industry analysts validate this pain.
- **Why Now?:** The thesis that the market is ripe for disruption is sound. Key tailwinds include a high readiness for AI adoption within VC, systemic market inefficiencies, and a fragmented competitive landscape of point solutions. The maturity of the underlying AI technology stack makes a comprehensive solution like this more feasible to build than ever before.
- **Market Size (TAM, SAM, SOM):** The market sizing appears realistic.
    - **TAM ($350M):** This figure, representing the annual revenue if every global investor/accelerator used the product, is plausible. With thousands of firms worldwide, an average ACV in the tens of thousands would easily support this total market size.
    - **SAM ($160M):** This represents firms with the most acute need for the solution and is a reasonable subset of the TAM.
    - **SOM ($1.4M ARR):** Their 3-year goal is to capture ~200 firms for a $1.4M ARR. This is a specific, believable, and attainable medium-term target. The market is unquestionably large enough to support a venture-scale outcome.
      (Note: The deck has a typo on the market size slide listing the SOM as $2M but later financial projections correctly calculate it as $1.4M based on stated assumptions)
      .

4. Solution & Product (Score: 6/10)

- **Clarity & Value Proposition:** The value proposition is exceptionally clear: an end-to-end AI platform that automates sourcing, screening, and analysis to save VCs time and money. If executed, this is a 10x improvement over the status quo of manually "stitching

together fragmented workflows." However, the deck lacks any product mockups, demos, or architectural details, keeping the solution purely conceptual at this stage.
- **Defensibility & Moat:** The proposed long-term moat is a "Proprietary Data Flywheel," where the platform gets smarter with more use. This is a classic and powerful SaaS moat, but it does not exist yet. Their initial moat is speed and focus, which is weak. They risk being out-maneuvered by competitors who could build similar features.
- **Product Roadmap:** The roadmap is logical and phased appropriately. It prioritizes validation with design partners (Q3 2025) , building the V1 core product (Q4 2025 - Q1 2026) , and then iterating to a V2 based on paying customer feedback (Q2 2026). This demonstrates a sound product strategy.

5. Go-to-Market & Business Model (Score: 8/10)

- **Customer Acquisition:** The GTM is specific and credible for an early-stage B2B company. They plan to start with pilots with 3-5 diverse firms to build case studies and testimonials, convert these pilots into their first paying customers, and then scale through a public launch.
- **Business Model:** The model is a straightforward B2B SaaS subscription. The proposed price of ~$6,999 ACV is a smart entry point. It's low enough to be an easy "yes" for a small fund or accelerator experiencing the pain, while high enough to build a real business. It effectively positions them as providing more value than expensive legacy tools.
- **Unit Economics:** It is too early to analyze LTV/CAC. However, a ~$7k ACV with a high-retention SaaS model targeting a market with inefficient communication has the potential for very healthy unit economics. The key will be whether they can keep customer acquisition costs low through their pilot-to-conversion strategy and avoid the need for an expensive enterprise sales force early on.

6. Traction & Progress

- **Current State:** The company is at the validation/idea stage. Tangible progress is limited to interviews with over 10 VCs and securing one firm for their Design Partner Program. The quotes provided are qualitative validation of the problem but do not represent product traction.
- **Velocity:** The timeline indicates validation and fundraising prep occurred in Q3 2025 (the current quarter). Given this, the velocity is reasonable for a team that has just started, but there is no history of rapid execution to evaluate yet.

7. Competitive Landscape

- **Competitor Analysis:** The team has a good understanding of the competitive landscape, correctly identifying the different categories of tools VCs currently use: in-house AI (EQT's Motherbrain) , point solutions (Connectic's Wendal) , and integrated CRMs (Affinity, edda). They have not missed any obvious major players.

- **Competitive Advantage:** Their stated advantage is being the first "true intelligent platform" that offers a unified, end-to-end workflow rather than a fragmented point solution. This is a powerful strategic advantage *if* they can execute on the product vision. The challenge will be making each part of their "end-to-end" platform as good as the best-in-class point solution for that specific task.

8. Critical Risks & Mitigation

- **Execution Risk: (High)** This is the most significant risk. Can a two-person team with unknown backgrounds build and deliver a complex, end-to-end AI platform that VCs will trust with their core workflow? The deck does not provide any information to mitigate this risk.
- **Competitive Risk: (High)** What prevents a well-funded incumbent like Affinity or a data provider like PitchBook from building an "AI agent" feature on top of their existing products and user bases? AgentVC's moat is currently non-existent, making them vulnerable to fast followers. Their mitigation is to build a better, workflow-native product, which is an execution challenge.
- **Market Risk / "Good Enough" Problem: (Medium)** VCs have operated with inefficient, "stitched-together" workflows for decades. There's a risk that these existing solutions, while not optimal, are "good enough" and that the activation energy required to switch to and trust a new, unproven platform for such a critical function will be too high. Their pilot program is a good mitigation strategy to directly address this by proving value with early customers.

9. Key Questions for the Founders (Action Items)

1. Please walk us through your professional backgrounds. What specific experiences led you to identify this problem, and what gives you an "unfair advantage" in solving it?
2. Can you provide more detail on the technical architecture? What parts are you building from scratch versus leveraging existing APIs or models?
3. Your moat relies on a "proprietary data flywheel." How exactly does this flywheel work? What unique data will you be capturing that competitors cannot?
4. Your first design partner is a critical piece of validation. Can you share the specifics of this engagement? What are the key success metrics for the pilot, and what is the path to converting them to a paying customer?
5. How do you plan to handle the "cold start" problem? How will the AI provide valuable insights for the first customer before the data flywheel is spinning?
6. The ~$7k price point is accessible. How do you see ACV expanding over time? What does your future pricing model and product tiering look like?
7. Beyond sourcing and analysis, what is your long-term vision for the "end-to-end" platform? Does it eventually include portfolio management, reporting, or other VC functions?

10. Final Recommendation & Rationale

Final Call: WATCHLIST

Justification: AgentVC presents a compelling vision to solve a real, high-value problem for a lucrative customer base. Their strategic positioning as a unified, intelligent platform is spot-on, and their initial GTM and pricing models are sound. However, the investment is untenable at this stage due to a critical lack of information about the founding team. Without knowing their background, we cannot assess their ability to execute on the ambitious technical and business plan they've laid out. The single biggest reason to pass *for now* is the execution risk associated with an unknown team. I would become interested in investing if they can (1) successfully launch their V1 product with their initial design partners, and (2) convert those partners into their first paying customers, thus proving both product-market fit and their ability to execute.